

Foundations of Data Mining: Assignment 4

- Group: 20
- Students: Junzhou Jiang (0963164) / Jiyang Li (0975662)

In [10]:

```
%matplotlib inline
from preamble import *
plt.rcParams['savefig.dpi'] = 100 # This controls the size of your figures
# Comment out and restart notebook if you only want the last output of
each cell.
InteractiveShell.ast_node_interactivity = "all"
```

Backpropagation (3 points)

Solution



- We choose Loss function:

$$\blacksquare L(x, y; w) = \frac{1}{2}(o_{\theta} - y)^2$$

- ϕ is ReLU= $\max(q, 0)$ here
- Forward 1:

$$r = \phi(q_0) = \phi(w_0 x_0) = 1$$

$$o = \phi(q_1) = \phi(w_1 r) = 2$$

- Backward 1:

$$w_1 = w_1 - \alpha \frac{\partial L}{\partial w_1} = w_1 - \alpha \frac{\partial L}{\partial o} \frac{\partial o}{\partial q_1} \frac{\partial q_1}{\partial w_1} = w_1 - \alpha \cdot x_0(o - y_0) = 2 - 0.1 * 1 * 0 = 2$$

$$w_0 = w_0 - \alpha \frac{\partial L}{\partial w_0} = w_0 - \alpha \frac{\partial L}{\partial o} \frac{\partial o}{\partial w_0} = w_0 - 0 = 1$$

Since $o_{\theta} = y$ for first point, weights don't update.

- Forward 2:

$$r = \phi(q_0) = \phi(w_0 x_1) = 2$$

$$o = \phi(q_1) = \phi(w_1 r) = 4$$

- Backward 2:

$$\frac{\partial L}{\partial w_1} \quad \frac{\partial L}{\partial a} \frac{\partial o}{\partial a} \frac{\partial q_1}{\partial w_1}$$

$$w_1 = w_1 - \alpha \frac{\partial L}{\partial w_1} = w_1 - \alpha \frac{\partial L}{\partial q_1} \frac{\partial q_1}{\partial r} \frac{\partial r}{\partial q_0} \frac{\partial q_0}{\partial w_1} = w_1 - \alpha \cdot x_1(o - y_1) = 2 - 0.1 * 2 * 1 = 1.8$$

$$w_0 = w_0 - \alpha \frac{\partial L}{\partial w_0} = w_0 - \alpha \frac{\partial L}{\partial q_1} \frac{\partial q_1}{\partial r} \frac{\partial r}{\partial q_0} \frac{\partial q_0}{\partial w_0} = w_0 - \alpha * x_1 * w_1(o - y_1) = 0.6$$

$o_\theta \neq y$ for second point, so weights update after backpropa.

Training Deep Models (3 points)

The model in the example code below performs poorly as its depth increases. Train this model on the MNIST digit detection task.

Examine its training performance by gradually increasing its depth:

- Set the depth to 1 hidden layer
- Set the depth to 2 hidden layers
- Set the depth to 3 hidden layers

Modify the model such that you improve its performance when its depth increases. Train the new model again for the different depths:

- Set the depth to 1 hidden layer
- Set the depth to 2 hidden layers
- Set the depth to 3 hidden layers

Submit an explanation for the limitation of the original model. Explain your modification. Submit your code and 6 plots (can be overlaid) for the training performance of both models with different depths.

In [11]:

```
# (You don't need to change this part of the code)
from __future__ import print_function
import numpy as np
np.random.seed(1234)

from keras.datasets import mnist
from keras.models import Sequential
from keras.layers.core import Dense, Dropout, Activation
from keras.optimizers import SGD
from keras.utils import np_utils

import matplotlib.pyplot as plt

batch_size = 128
nb_classes = 10
nb_epoch = 10
```

In [12]:

```
# (You don't need to change this part of the code)
# the data, shuffled and split between train and test sets
(X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train.reshape(60000, 784)
X_test = X_test.reshape(10000, 784)
```

```

X_train = X_train.astype('float32')
X_test = X_test.astype('float32')
X_train /= 255
X_test /= 255
print(X_train.shape[0], 'train samples')
print(X_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
Y_train = np_utils.to_categorical(y_train, nb_classes)
Y_test = np_utils.to_categorical(y_test, nb_classes)

```

60000 train samples
10000 test samples

In [13]:

```

# Use this parameter to change the depth of the model
number_hidden_layers = 3 # Number of hidden layers

```

In [4]:

```

def train_model(number_hidden_layers, activation):
    model = Sequential()
    model.add(Dense(512, input_shape=(784,), activation=activation))
    model.add(Dropout(0.2))

    while number_hidden_layers > 1:
        model.add(Dense(512))
        model.add(Activation(activation))
        model.add(Dropout(0.2))
        number_hidden_layers -= 1

    model.add(Dense(10))
    model.add(Activation('softmax'))
    model.summary()

    model.compile(loss='categorical_crossentropy',
                  optimizer=SGD(),
                  metrics=['accuracy'])

    history = model.fit(X_train, Y_train,
                        batch_size=batch_size, nb_epoch=nb_epoch,
                        verbose=0, validation_data=(X_test, Y_test))
    score = model.evaluate(X_test, Y_test, verbose=0)

    return score[0], score[1], history

def run_model(activation):
    for num_layer in [1, 2, 3]:
        test_score, test_accuracy, history =
train_model(num_layer, activation)
        print('Depth:', num_layer)
        print('Test score:', test_score)
        print('Test accuracy:', test_accuracy)

        plt.subplot(1, 3, num_layer)
        plt.plot(history.history['acc'])
        plt.plot(history.history['val_acc'])
        plt.title('model accuracy with depth: %d' % (num_layer))

```

```
plt.ylabel('accuracy')
plt.xlabel('epoch')
plt.legend(['train', 'test'], loc='upper left')
```

In [5]:

```
plt.figure(figsize=(12, 4))
plt.suptitle('Sigmoid Activation')
run_model('sigmoid')
plt.show()
```

Out[5]:

<matplotlib.figure.Figure at 0x116b693c8>

Out[5]:

<matplotlib.text.Text at 0x116b3f940>

Layer (type)	Output Shape	Param #
dense_1 (Dense)	(None, 512)	401920
dropout_1 (Dropout)	(None, 512)	0
dense_2 (Dense)	(None, 10)	5130
activation_1 (Activation)	(None, 10)	0

=====
Total params: 407,050.0
Trainable params: 407,050.0
Non-trainable params: 0.0

Depth: 1
Test score: 0.466248050857
Test accuracy: 0.8811

Layer (type)	Output Shape	Param #
dense_3 (Dense)	(None, 512)	401920
dropout_2 (Dropout)	(None, 512)	0
dense_4 (Dense)	(None, 512)	262656
activation_2 (Activation)	(None, 512)	0
dropout_3 (Dropout)	(None, 512)	0
dense_5 (Dense)	(None, 10)	5130
activation_3 (Activation)	(None, 10)	0

=====
Total params: 669,706.0
Trainable params: 669,706.0
Non-trainable params: 0.0

Depth: 2
Test score: 0.869930342674
Test accuracy: 0.7908

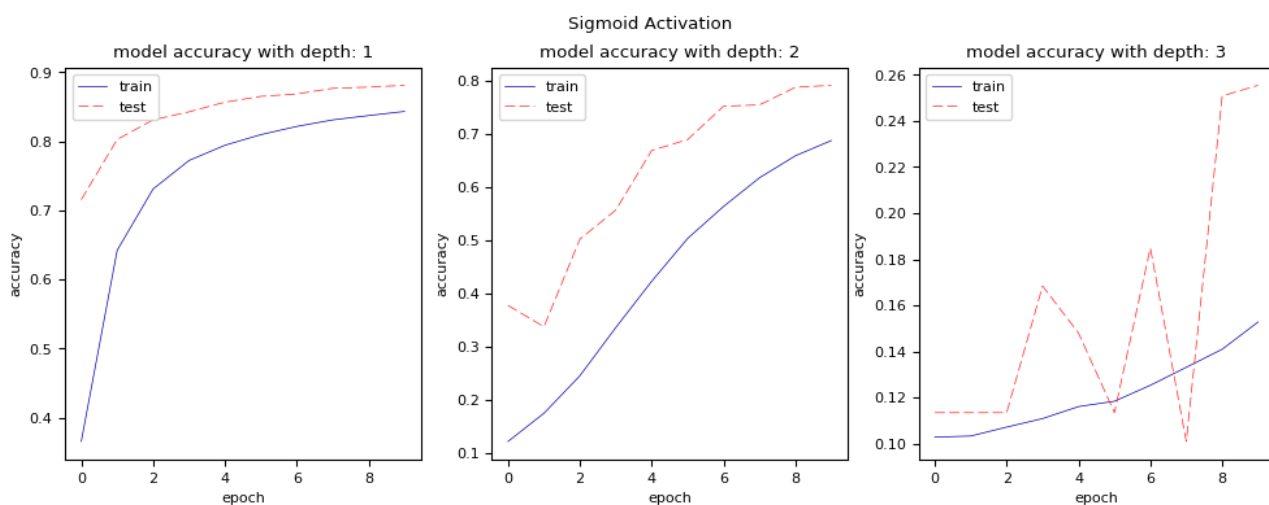
Layer (type)	Output Shape	Param #
--------------	--------------	---------

layer (type)	Output Shape	Param #
dense_6 (Dense)	(None, 512)	401920
dropout_4 (Dropout)	(None, 512)	0
dense_7 (Dense)	(None, 512)	262656
activation_4 (Activation)	(None, 512)	0
dropout_5 (Dropout)	(None, 512)	0
dense_8 (Dense)	(None, 512)	262656
activation_5 (Activation)	(None, 512)	0
dropout_6 (Dropout)	(None, 512)	0
dense_9 (Dense)	(None, 10)	5130
activation_6 (Activation)	(None, 10)	0
=====		
Total params: 932,362.0		
Trainable params: 932,362.0		
Non-trainable params: 0.0		

Depth: 3

Test score: 2.23485740509

Test accuracy: 0.2554



In [20]:

```
plt.figure(figsize=(12, 4))
plt.suptitle('ReLU Activation')
run_model('relu')
plt.show()
```

Out[20]:

<matplotlib.figure.Figure at 0x1316d2518>

Out[20]:

<matplotlib.text.Text at 0x1316d2358>

Layer (type)	Output Shape	Param #
--------------	--------------	---------

```

=====
dense_30 (Dense)                (None, 512)                401920
-----
dropout_20 (Dropout)            (None, 512)                0
-----
dense_31 (Dense)                (None, 10)                 5130
-----
activation_20 (Activation)      (None, 10)                 0
=====
Total params: 407,050.0
Trainable params: 407,050.0
Non-trainable params: 0.0

```

```

Depth: 1
Test score: 0.256563024029
Test accuracy: 0.9293

```

Layer (type)	Output Shape	Param #
dense_32 (Dense)	(None, 512)	401920
dropout_21 (Dropout)	(None, 512)	0
dense_33 (Dense)	(None, 512)	262656
activation_21 (Activation)	(None, 512)	0
dropout_22 (Dropout)	(None, 512)	0
dense_34 (Dense)	(None, 10)	5130
activation_22 (Activation)	(None, 10)	0

```

=====
Total params: 669,706.0
Trainable params: 669,706.0
Non-trainable params: 0.0

```

```

Depth: 2
Test score: 0.207567299627
Test accuracy: 0.941

```

Layer (type)	Output Shape	Param #
dense_35 (Dense)	(None, 512)	401920
dropout_23 (Dropout)	(None, 512)	0
dense_36 (Dense)	(None, 512)	262656
activation_23 (Activation)	(None, 512)	0
dropout_24 (Dropout)	(None, 512)	0
dense_37 (Dense)	(None, 512)	262656
activation_24 (Activation)	(None, 512)	0
dropout_25 (Dropout)	(None, 512)	0
dense_38 (Dense)	(None, 10)	5130

```

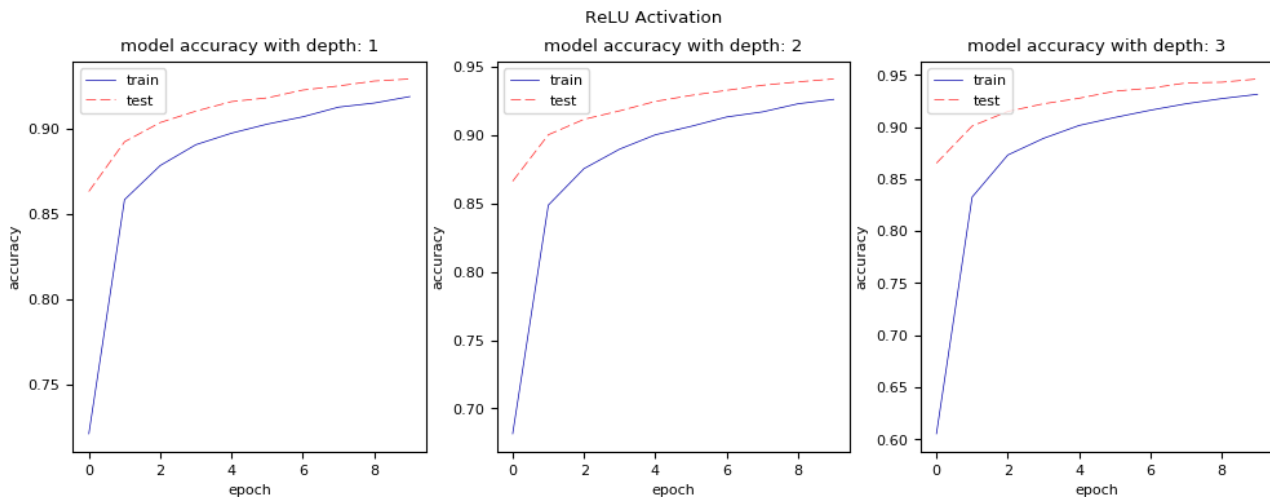
activation_25 (Activation)      (None, 10)                      0
=====
Total params: 932,362.0
Trainable params: 932,362.0
Non-trainable params: 0.0

```

```

Depth: 3
Test score: 0.180363421577
Test accuracy: 0.946

```



Solution

We can conclude from figures above that changing the activation 'sigmoid' to 'ReLU' can modify the results a lot.

Sigmoid activation function is $\frac{1}{1+e^{-w^T x}}$, so the gradient calculation of it would be complex. Although it seems good when number of hidden_layer is 1 but when we add more layers, the gradient would vanish or explode thus makes little or too heavy contribution to the network and result in low and unstable accuracy.

On the contrary, the ReLU is a max function: $\max\{0, x\}$, so we only have gradient 0 or 1, which is simple to compute. And without the gradient vanishing and exploding problems, the accuracy will be higher with the number of layers increasing.

Convolutional Neural Networks for Filtering (2 points)

Solution

We define the number of layer as 3. Which includes the input layer, the sepia combine Gaussian blurring layer and the output layer.

The image consists of 128×128 pixels. The color of each point is defined by a 3D vector, and the value of each axis of the vector represents the RGB value respectively. The middle layer is achieved by a $3 \times \text{Conv.} \times 5 \times 5 \times 3$ convolution. The neuron computes the RGB values respectively with both sepia effect and Gaussian blurring effect.

As a discretized version of the filter is given, which can be represent by a matrix like below:

$$B'(i, j) = \sum_{m=1}^5 \sum_{n=1}^5 ((R_{(i+m-1, j+n-1)} * .272) + (G_{(i+m-1, j+n-1)} * .534) + (B_{(i+m-1, j+n-1)} * .131)) \times Gus(m, i, j \text{ is from } 1 \text{ to } 124).$$

To avoid overfitting we use dropout to regularize the model. With each state h_t decided by h_{t-1} and h_t^{l-1} , where t represents the time and l represents the layer, we choose to apply dropout on h_t^{l-1} part, which referred from [RECURRENT NEURAL NETWORK REGULARIZATION](#).

- Input data format: a short video clip consisting several frames.
- Number of layers: 4
- Type of layer: Input layer (3D), Convolutional layer (3D), Recurrent layer (2 LSTM layers each is 1D), and Output layer (1D vector).
- Regularization: dropout
- loss function: cross entropy

MNIST Calculator (5 points)

In [35]:

```
from __future__ import print_function
from keras.models import Sequential
from keras.engine.training import _slice_arrays
from keras.layers import Activation, TimeDistributed, Dense, RepeatVector,
recurrent
import numpy as np
from six.moves import range
from keras.datasets import mnist
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

DIGITS = 3
MAXLEN = DIGITS + 1 + DIGITS
chars = '0123456789+ '
RNN = recurrent.LSTM
HIDDEN_SIZE = 256
BATCH_SIZE = 512
LAYERS = 2
TRAINING_SIZE = 50000

class CharacterTable(object):
    def __init__(self, chars, maxlen):
        self.chars = sorted(set(chars))
        self.char_indices = dict((c, i) for i, c in enumerate(self.chars))
        self.indices_char = dict((i, c) for i, c in enumerate(self.chars))
        self.maxlen = maxlen

    def encode(self, C, maxlen=None):
        maxlen = maxlen if maxlen else self.maxlen
        X = np.zeros((maxlen, len(self.chars)))
        for i, c in enumerate(C):
            X[i, self.char_indices[c]] = 1
        return X

    def decode(self, X, calc_argmax=True):
        if calc_argmax:
            X = X.argmax(axis=-1)
        return ''.join(self.indices_char[x] for x in X)

class colors:
```

```
ok = '\033[92m'
fail = '\033[91m'
close = '\033[0m'
```

In [36]:

```
ctable = CharacterTable(chars, MAXLEN)
```

In [37]:

```
img_rows = img_cols = 28

(x_train, y_train), (x_test, y_test) = mnist.load_data()

if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)
y_train1 = y_train
```

In [40]:

```
def random_question():
    f = lambda: int(''.join(np.random.choice(list('0123456789')) for i in range(np.random.randint(1, DIGITS + 1))))
    a, b = f(), f()

    query = '{}+{}'.format(a, b)
    ans = str(a + b)

    query += ' ' * (MAXLEN - len(query))
    ans += ' ' * (DIGITS + 1 - len(ans))

    return query, ans

def encode_question(question, output=None):
    if output is None:
        i = 0
        output = np.zeros([1, MAXLEN] + list(input_shape), dtype=x_train.dtype)

        for j, char in enumerate(question):
            if char in '0123456789':
                orin = None
                while orin is None:
                    num = np.random.choice(y_train1.shape[0])
                    if y_train1[num] == int(char):
                        orin = x_train[num]
            else:
                orin = 255 * np.random.random(input_shape)

            output[i, j] = orin
    else:
        for j, char in enumerate(question):
            if char in '0123456789':
                orin = None
                while orin is None:
```

```

        num = np.random.choice(y_train1.shape[0])
        if y_train1[num] == int(char):
            orin = x_train[num]
        else:
            orin = 255 * np.random.random(input_shape)

        output[j] = orin
    return output

```

In [41]:

```

X = np.zeros([TRAINING_SIZE, MAXLEN] + list(input_shape), dtype=x_train.dtype)
y = np.zeros((TRAINING_SIZE, DIGITS + 1, len(chars)), dtype=np.bool)

for i in range(TRAINING_SIZE):
    question, answer = random_question()
    y[i] = ctable.encode(answer, maxlen=DIGITS + 1)
    X[i] = encode_question(question, X[i])

# Shuffle (X, y) in unison as the later parts of X will almost all be large r digits
indices = np.arange(len(y))
np.random.shuffle(indices)
X = X[indices]
y = y[indices]

split_at = int(X.shape[0] * 0.9)
(X_train, X_val) = (_slice_arrays(X, 0, split_at), _slice_arrays(X, split_at))
(y_train, y_val) = (y[:split_at], y[split_at:])

```

In [42]:

```

model = Sequential()
model.add(TimeDistributed(Conv2D(32, kernel_size=(5, 5),
                                activation='relu', strides=(3, 3)),
                                input_shape=X_train.shape[1:])))
model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))
model.add(Dropout(0.25))
model.add(TimeDistributed(Conv2D(64, (3, 3), activation='relu', strides=(1, 1)
                                )))
model.add(TimeDistributed(MaxPooling2D(pool_size=(2, 2))))
model.add(Dropout(0.5))
model.add(TimeDistributed(Flatten()))
model.add(TimeDistributed(Dense(128, activation='relu')))

model.add(RNN(HIDDEN_SIZE))
model.add(RepeatVector(DIGITS + 1))
for _ in range(LAYERS):
    model.add(RNN(HIDDEN_SIZE, return_sequences=True))

# For each of step of the output sequence, decide which character should be chosen
model.add(TimeDistributed(Dense(len(chars))))
model.add(Activation('softmax'))

model.compile(loss='categorical_crossentropy',

```

```
optimizer='adam',  
metrics=['accuracy'])
```

In [43]:

```
for iteration in range(1, 10):  
    print()  
    print('-' * 50)  
    print('Iteration', iteration)  
    model.fit(X_train, y_train, validation_data=(X_val, y_val),  
              batch_size=BATCH_SIZE, epochs=5, verbose=2)  
  
    for i in range(10):  
        question, answer = random_question()  
        out = encode_question(question)  
        preds = model.predict_classes(out, verbose=0)  
        guess = ctable.decode(preds[0], calc_argmax=False)  
        print('Q', question)  
        print('T', answer)  
        print(colors.ok + '☑' + colors.close if answer == guess else colors.  
fail + '☒' + colors.close, guess)  
        print('----')
```

```
-----  
Iteration 1  
Train on 45000 samples, validate on 5000 samples  
Epoch 1/5  
111s - loss: 1.7623 - acc: 0.4094 - val_loss: 1.5621 - val_acc: 0.4361  
Epoch 2/5  
108s - loss: 1.5418 - acc: 0.4505 - val_loss: 1.5156 - val_acc: 0.4544  
Epoch 3/5  
113s - loss: 1.5180 - acc: 0.4525 - val_loss: 1.4950 - val_acc: 0.4581  
Epoch 4/5  
114s - loss: 1.5008 - acc: 0.4571 - val_loss: 1.4631 - val_acc: 0.4666  
Epoch 5/5  
117s - loss: 1.4619 - acc: 0.4663 - val_loss: 1.3843 - val_acc: 0.4885
```

Out[43]:

```
<keras.callbacks.History at 0x144953550>
```

```
Q 8+9  
T 17  
☒ 13  
---  
Q 2+69  
T 71  
☒ 68  
---  
Q 500+53  
T 553  
☒ 558  
---  
Q 9+4  
T 13  
☒ 15  
---  
Q 778+77  
T 855  
☒ 908  
---
```

Q 114+877

T 991

☒ 111

Q 99+12

T 111

☒ 12

Q 6+4

T 10

☒ 1

Q 5+995

T 1000

☒ 908

Q 2+4

T 6

☒ 1

Iteration 2

Train on 45000 samples, validate on 5000 samples

Epoch 1/5

112s - loss: 1.3808 - acc: 0.4906 - val_loss: 1.2925 - val_acc: 0.5095

Epoch 2/5

109s - loss: 1.3092 - acc: 0.5132 - val_loss: 1.1946 - val_acc: 0.5571

Epoch 3/5

109s - loss: 1.1999 - acc: 0.5530 - val_loss: 1.0114 - val_acc: 0.6270

Epoch 4/5

108s - loss: 1.0816 - acc: 0.6063 - val_loss: 0.8796 - val_acc: 0.6932

Epoch 5/5

109s - loss: 0.9815 - acc: 0.6535 - val_loss: 0.7769 - val_acc: 0.7422

Out[43]:

<keras.callbacks.History at 0x14544d710>

Q 36+5

T 41

☒ 50

Q 952+69

T 1021

☒ 1011

Q 104+4

T 108

☒ 118

Q 412+423

T 835

☒ 645

Q 29+7

T 36

☒ 36

Q 0+48

T 48

```

❌ 58
---
Q 6+8
T 14
✅ 14
---
Q 534+0
T 534
❌ 434
---
Q 2+40
T 42
✅ 42
---
Q 98+12
T 110
❌ 100
---

-----
Iteration 3
Train on 45000 samples, validate on 5000 samples
Epoch 1/5
113s - loss: 0.8899 - acc: 0.7014 - val_loss: 0.6701 - val_acc: 0.7941
Epoch 2/5
757s - loss: 0.7994 - acc: 0.7408 - val_loss: 0.5679 - val_acc: 0.8357
Epoch 3/5
113s - loss: 0.7312 - acc: 0.7680 - val_loss: 0.5123 - val_acc: 0.8515
Epoch 4/5
115s - loss: 0.6747 - acc: 0.7899 - val_loss: 0.4494 - val_acc: 0.8702
Epoch 5/5
114s - loss: 0.6237 - acc: 0.8073 - val_loss: 0.4324 - val_acc: 0.8698

Out[43]:
<keras.callbacks.History at 0x14548c668>

Q 18+53
T 71
✅ 71
---
Q 82+78
T 160
✅ 160
---
Q 11+2
T 13
✅ 13
---
Q 435+69
T 504
❌ 404
---
Q 2+1
T 3
❌ 2
---
Q 73+2
T 75
✅ 75
---

```

Q 48+194

T 842

✖ 830

Q 31+11

T 42

✔ 42

Q 169+543

T 712

✖ 8022

Q 943+392

T 1335

✖ 1315

Iteration 4

Train on 45000 samples, validate on 5000 samples

Epoch 1/5

115s - loss: 0.5910 - acc: 0.8202 - val_loss: 0.3852 - val_acc: 0.8902

Epoch 2/5

117s - loss: 0.5602 - acc: 0.8309 - val_loss: 0.3631 - val_acc: 0.8930

Epoch 3/5

121s - loss: 0.5280 - acc: 0.8418 - val_loss: 0.3308 - val_acc: 0.9056

Epoch 4/5

125s - loss: 0.5058 - acc: 0.8493 - val_loss: 0.3167 - val_acc: 0.9117

Epoch 5/5

126s - loss: 0.4895 - acc: 0.8552 - val_loss: 0.3136 - val_acc: 0.9130

Out[43]:

<keras.callbacks.History at 0x14548c710>

Q 8+20

T 28

✔ 28

Q 266+493

T 759

✖ 749

Q 478+66

T 544

✔ 544

Q 3+763

T 766

✔ 766

Q 932+74

T 1006

✖ 100

Q 73+15

T 88

✔ 88

Q 97+187

T 284

✖ 254

Q 4+648

T 652

Q 50+611

T 661

Q 991+9

T 1000

Q 53+626

T 679

Q 599+1

T 600

Q 5+3

T 8

Q 25+956

T 981

Q 7+78

T 85

Q 44+90

T 134

Q 381+55

T 436

Q 15+91

T 652

T 661

T 1000

T 679

T 600

T 8

T 981

T 85

T 134

T 436

T 652

T 661

T 1000

T 679

T 600

T 8

T 981

T 85

T 134

T 436

T 652

T 661

T 1000

T 679

T 600

T 8

T 981

T 85

T 134

T 436

T 652

T 661

T 1000

T 679

T 600

T 8

T 981

T 85

T 134

T 436

T 652

T 661

T 1000

T 679

T 600

T 8

T 981

T 85

T 134

T 436

T 652

T 661

T 1000

T 679


```
T 106
✓ 106
---
Q 0+91
T 91
✓ 91
---
Q 7+6
T 13
✓ 13
---
```

```
-----
Iteration 6
Train on 45000 samples, validate on 5000 samples
Epoch 1/5
118s - loss: 0.4133 - acc: 0.8814 - val_loss: 0.2539 - val_acc: 0.9301
Epoch 2/5
120s - loss: 0.4051 - acc: 0.8837 - val_loss: 0.2462 - val_acc: 0.9310
Epoch 3/5
116s - loss: 0.3953 - acc: 0.8873 - val_loss: 0.2383 - val_acc: 0.9353
Epoch 4/5
117s - loss: 0.3881 - acc: 0.8892 - val_loss: 0.2441 - val_acc: 0.9331
Epoch 5/5
125s - loss: 0.3840 - acc: 0.8902 - val_loss: 0.2290 - val_acc: 0.9357
```

Out[43]:

```
<keras.callbacks.History at 0x13c6ba320>
```

```
Q 12+587
T 599
✗ 699
---
Q 3+8
T 11
✓ 11
---
Q 934+794
T 1728
✓ 1728
---
Q 0+7
T 7
✓ 7
---
Q 1+6
T 7
✓ 7
---
Q 50+54
T 104
✓ 104
---
Q 42+922
T 964
✓ 964
---
Q 719+59
T 778
✓ 778
```

```
---
Q 377+155
T 532
☒ 432
---
```

```
Q 78+5
T 83
☒ 13
---
```

```
-----
Iteration 7
Train on 45000 samples, validate on 5000 samples
Epoch 1/5
119s - loss: 0.3750 - acc: 0.8936 - val_loss: 0.2273 - val_acc: 0.9394
Epoch 2/5
117s - loss: 0.3733 - acc: 0.8946 - val_loss: 0.2287 - val_acc: 0.9368
Epoch 3/5
115s - loss: 0.3637 - acc: 0.8969 - val_loss: 0.2147 - val_acc: 0.9407
Epoch 4/5
115s - loss: 0.3568 - acc: 0.8993 - val_loss: 0.2130 - val_acc: 0.9438
Epoch 5/5
115s - loss: 0.3544 - acc: 0.8995 - val_loss: 0.2221 - val_acc: 0.9406
```

Out[43]:

<keras.callbacks.History at 0x14548ca58>

```
Q 6+7
T 13
☑ 13
---
```

```
Q 667+6
T 673
☑ 673
---
```

```
Q 787+17
T 804
☒ 704
---
```

```
Q 8+62
T 70
☑ 70
---
```

```
Q 625+6
T 631
☑ 631
---
```

```
Q 35+0
T 35
☑ 35
---
```

```
Q 357+9
T 366
☒ 361
---
```

```
Q 86+748
T 834
☑ 834
---
```

```
Q 98+251
T 349
```

```
- - -
✓ 349
---
Q 3+300
T 303
✓ 303
---

-----
Iteration 8
Train on 45000 samples, validate on 5000 samples
Epoch 1/5
119s - loss: 0.3478 - acc: 0.9019 - val_loss: 0.2143 - val_acc: 0.9437
Epoch 2/5
119s - loss: 0.3396 - acc: 0.9044 - val_loss: 0.2142 - val_acc: 0.9439
Epoch 3/5
126s - loss: 0.3418 - acc: 0.9032 - val_loss: 0.2033 - val_acc: 0.9456
Epoch 4/5
122s - loss: 0.3320 - acc: 0.9059 - val_loss: 0.2095 - val_acc: 0.9429
Epoch 5/5
117s - loss: 0.3329 - acc: 0.9050 - val_loss: 0.1952 - val_acc: 0.9467

Out[43]:

<keras.callbacks.History at 0x14548c5c0>

Q 55+8
T 63
✓ 63
---
Q 1+485
T 486
✓ 486
---
Q 578+53
T 631
✓ 631
---
Q 1+888
T 889
✓ 889
---
Q 18+308
T 326
✓ 326
---
Q 2+971
T 973
✓ 973
---
Q 927+55
T 982
✓ 982
---
Q 89+3
T 92
✗ 83
---
Q 995+80
T 1075
✗ 1065
---
```

Q 9+84

T 93

✓ 93

Iteration 9

Train on 45000 samples, validate on 5000 samples

Epoch 1/5

123s - loss: 0.3280 - acc: 0.9080 - val_loss: 0.1959 - val_acc: 0.9484

Epoch 2/5

117s - loss: 0.3259 - acc: 0.9083 - val_loss: 0.1978 - val_acc: 0.9467

Epoch 3/5

121s - loss: 0.3224 - acc: 0.9084 - val_loss: 0.2238 - val_acc: 0.9414

Epoch 4/5

123s - loss: 0.3196 - acc: 0.9096 - val_loss: 0.2033 - val_acc: 0.9468

Epoch 5/5

125s - loss: 0.3151 - acc: 0.9114 - val_loss: 0.1930 - val_acc: 0.9482

Out[43]:

<keras.callbacks.History at 0x14548c9e8>

Q 89+92

T 181

✓ 181

Q 92+91

T 183

✓ 183

Q 7+24

T 31

✓ 31

Q 97+1

T 98

✓ 98

Q 87+83

T 170

✓ 170

Q 43+28

T 71

✗ 62

Q 89+62

T 151

✓ 151

Q 6+39

T 45

✓ 45

Q 934+760

T 1694

✓ 1694

Q 1+3

T 4

✓ 4

✓ 4

In []: